

STAT 341: Assignment 2

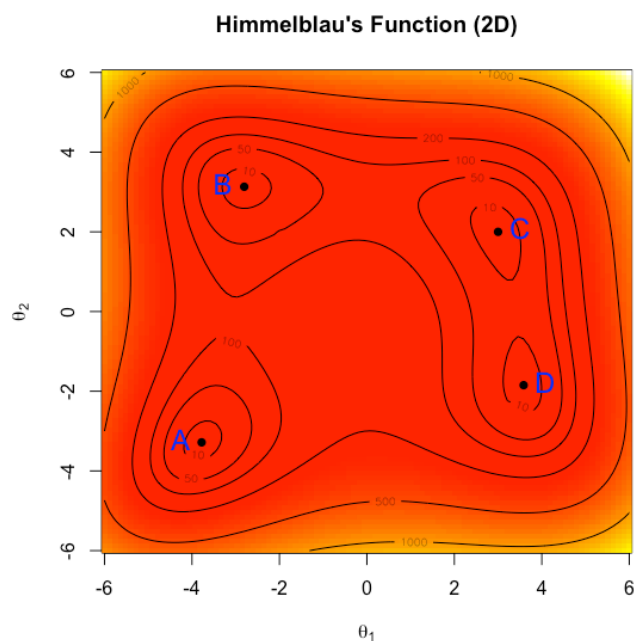

DUE: Friday February 14 by 11:59pm EST

QUESTION 1: Himmelblau's Function [18 points]

Himmelblau's Function is one of many non-convex [test functions](#) used for evaluating the performance of optimization methods. For $\theta \in \mathbb{R}^2$ the function is defined as follows

$$\rho(\theta) = (\theta_1^2 + \theta_2 - 11)^2 + (\theta_1 + \theta_2^2 - 7)^2$$

The figures below depict the function (as a 3-dimensional surface and with 2-dimensional contours) for $\theta_1 \in [-6, 6]$ and $\theta_2 \in [-6, 6]$.



As can be seen on the surfaces, this function has 4 local minima labelled (in the contour plot) $A = (-3.779310, -3.283186)$, $B = (-2.805118, 3.131312)$, $C = (3, 2)$, $D = (3.584428, -1.848126)$. Himmelblau's Function attains a value of zero at each of these four points.

- (a) [4 points] Write **rho** and **gradient** functions for Himmelblau's Function which take a single vector-valued input **theta**. Note that you may use without proof or derivation the fact that

$$\frac{\partial \rho}{\partial \theta_1} = 4\theta_1^3 + 4\theta_1\theta_2 - 42\theta_1 + 2\theta_2^2 - 14$$

and

$$\frac{\partial \rho}{\partial \theta_2} = 2\theta_1^2 + 4\theta_1\theta_2 + 4\theta_2^3 - 26\theta_2 - 22$$

Solution:

```
rho <- function(theta){
  (theta[1]^2 + theta[2] - 11)^2 + (theta[1] + theta[2]^2 - 7)^2
}

gradient <- function(theta){
  c(4*theta[1]^3 + 4*theta[1]*theta[2] - 42*theta[1] + 2*theta[2]^2 - 14,
    2*theta[1]^2 + 4*theta[1]*theta[2] + 4*theta[2]^3 - 26*theta[2] - 22)
}
```

- (b) [5 points] In this question you will explore the surface of the Himmelblau Function using gradient descent. In particular you will consider 5 different starting values and explore the impact of changing one's starting location. Using the `gradientDescent` function (from class) together with the `gridLineSearch` and `testConvergence` functions (from class) as well as the `rho` and `gradient` functions from part (a), find the solution to

$$\underset{\theta \in \mathbb{R}^2}{\operatorname{argmin}} \rho(\theta)$$

for each of the following five starting values. In each case state which minima you've converged to (A, B, C, or D) and be sure to include the output from the `gradientDescent` function.

- i. $\hat{\theta} = (0, 0)$

Solution:

```
gradientDescent(theta = c(0,0), rhoFn = rho, gradientFn = gradient,
  lineSearchFn = gridLineSearch,
  testConvergenceFn = testConvergence)
```

```
## $theta
## [1] 3.000613 1.997116
##
## $converged
## [1] TRUE
##
## $iteration
## [1] 14
##
## $fnValue
## [1] 0.0001197231
```

Here we converged to C.

- ii. $\hat{\theta} = (0, 3)$

Solution:

```
gradientDescent(theta = c(0,3), rhoFn = rho, gradientFn = gradient,
  lineSearchFn = gridLineSearch,
  testConvergenceFn = testConvergence)
```

```
## $theta
## [1] -2.801746 3.130191
##
## $converged
## [1] TRUE
##
## $iteration
## [1] 8
```

```
##  
## $fnValue  
## [1] 0.0004144257
```

Here we converged to B.

iii. $\hat{\theta} = (3, 0)$

Solution:

```
gradientDescent(theta = c(3,0), rhoFn = rho, gradientFn = gradient,  
                 lineSearchFn = gridLineSearch,  
                 testConvergenceFn = testConvergence)
```

```
## $theta  
## [1] 3.580904 -1.846533  
##  
## $converged  
## [1] TRUE  
##  
## $iteration  
## [1] 15  
##  
## $fnValue  
## [1] 0.0006483197
```

Here we converged to D.

iv. $\hat{\theta} = (0, -3)$

Solution:

```
gradientDescent(theta = c(0,-3), rhoFn = rho, gradientFn = gradient,  
                 lineSearchFn = gridLineSearch,  
                 testConvergenceFn = testConvergence)
```

```
## $theta  
## [1] 3.579644 -1.842832  
##  
## $converged  
## [1] TRUE  
##  
## $iteration  
## [1] 14  
##  
## $fnValue  
## [1] 0.001431711
```

Here we converged to D.

v. $\hat{\theta} = (-3, 0)$

Solution:

```
gradientDescent(theta = c(-3,0), rhoFn = rho, gradientFn = gradient,  
                 lineSearchFn = gridLineSearch,  
                 testConvergenceFn = testConvergence)
```

```
## $theta  
## [1] -2.803536 3.134450  
##
```

```
## $converged
## [1] TRUE
##
## $iteration
## [1] 9
##
## $fnValue
## [1] 0.0004841347
```

Here we converged to B.

- (c) [5 points] Recreate the contour plot shown above. You may find the functions `outer`, `image`, and `contour` useful for this task. Include on this plot **green** triangles at each of the starting points specified in (b) as well as **green** line **segments** connecting these starting points with their respective points of convergence.

Solution:


```
rho <- function(theta1,theta2){
  (theta1^2 + theta2 - 11)^2 + (theta1 + theta2^2 - 7)^2
}
theta1 <- seq(-6,6,length=100)
theta2 <- seq(-6,6,length=100)
Rho <- outer(theta1,theta2,"rho")
image(theta1, theta2, Rho, col = heat.colors(1000) , xlab = bquote(theta[1]), ylab = bquote(theta[2]),
  main = "Himmelblau's Function (2D)")
contour(theta1,theta2,Rho,add=T, levels = c(10, 50, 100, 200, 500, 1000))
points(x = c(3, -2.805118, -3.779310, 3.584428), y = c(2, 3.131312, -3.283186, -1.848126),
  pch = 16)

A <- c(-3.779310, -3.283186)
B <- c(-2.805118, 3.131312)
C <- c(3,2)
D <- c(3.584428, -1.848126)

text(x = A[1], y = A[2], labels = "A", pos = 2, cex = 2, col = "blue")
text(x = B[1], y = B[2], labels = "B", pos = 2, cex = 2, col = "blue")
text(x = C[1], y = C[2], labels = "C", pos = 4, cex = 2, col = "blue")
text(x = D[1], y = D[2], labels = "D", pos = 4, cex = 2, col = "blue")

points(x = c(0,-3,0,3,0), y = c(0,0,3,0,-3), col = "green", pch = 17, cex = 2)
segments(x0 = 0, y0 = 0, x1 = C[1], y1 = C[2], col = "green", lwd = 2)
segments(x0 = 0, y0 = 3, x1 = B[1], y1 = B[2], col = "green", lwd = 2)
segments(x0 = 3, y0 = 0, x1 = D[1], y1 = D[2], col = "green", lwd = 2)
segments(x0 = 0, y0 = -3, x1 = D[1], y1 = D[2], col = "green", lwd = 2)
segments(x0 = -3, y0 = 0, x1 = B[1], y1 = B[2], col = "green", lwd = 2)
```

Himmelblau's Function (2D)



- (d) [2 points] Based on your investigations in parts (b) and (c) explain the importance of the starting value when performing non-convex optimization (when locating a global optimum is desired).

Solution: What is made clear in the previous questions is that different starting points can very easily lead to different local minima. And in a surface with a global minimum, if our starting point is very far away, it is highly likely that we won't find it.

- (e) [2 points] Using the `gradientDescent` function (from class) together with the `fixedLineSearch` and `testConvergence` functions (from class) as well as the `rho` and `gradient` functions from part (a), attempt to find the solution to

$$\operatorname{argmin}_{\theta \in \mathbb{R}^2} \rho(\theta)$$

starting from $\hat{\theta} = (0,0)$. Can you find input settings that lead to convergence? If so state them. Comment on the value of algorithmically choosing a step size versus using a fixed step size.

Solution: I could not find input values that led to convergence in a reasonable amount of time. This is the benefit of algorithmically chosen step sizes; they balance accurate convergence with quick convergence.

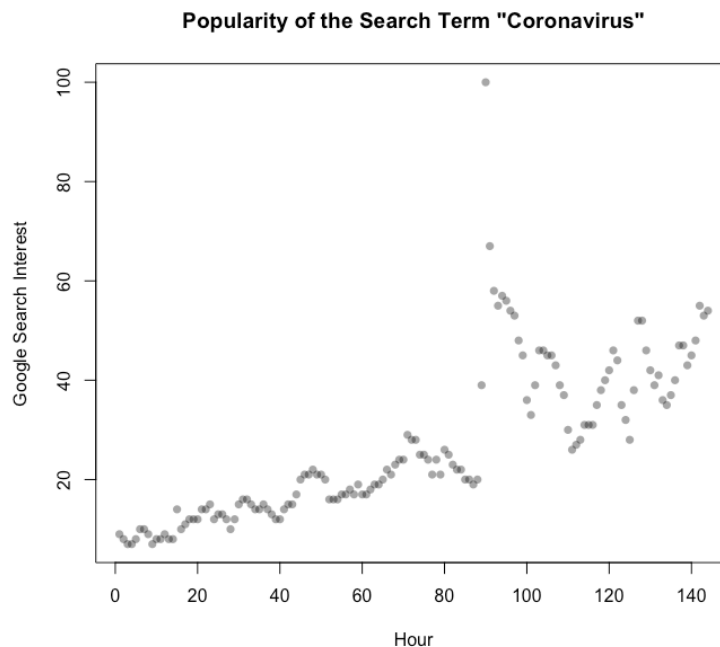
QUESTION 2: Google Trends “coronavirus” [25 points]

The 2019 Novel Coronavirus (2019-nCoV) is a recent strain of the coronavirus that impairs the respiratory systems of its host and that was first identified in Wuhan, Hubei Province, China. The first cases

were identified in December of 2019 and since then thousands of cases have been confirmed around the world. Unsurprisingly, amidst this outbreak citizens of the world have been turning to the internet for more information.

In this question we will analyze [Google Trends](#) data which illustrates the interest among Canadians in the search term “coronavirus”. The data available to you is stored in the `coronavirus.csv` file and summarized and visualized below. For each of the $N = 144$ hours in the time frame 1:00am January 21 – 1:00am January 28 we have recordings on the following variates:

Variate	Description
TimeStamp	A character string recorded as YYYY-MM-DDTHH indicating hours (i.e., 2020-01-22T05 corresponds to 5:00-6:00am on January 22nd)
h	A numeric variate indicating hours (i.e., $h=1$ corresponds to 1:00-2:00am on January 21st and so on).
interest	A numeric variate that represents search interest. A value of 100 is the peak popularity for the term. A value of 50 means that the term is half as popular, etc.



- (a) [2 points] As is demonstrated in the plot above, there is one hour in which Google searches for “coronavirus” in Canada increased dramatically. Determine the `TimeStamp` associated with this largest `interest` value. What is significant about this hour? (Hint: think current events).

Solution:

```
virus <- read.csv("/Users/nstevens/Dropbox/Teaching/STAT_341/Assignments/Assignment2/coronavirus.csv",
                 header = TRUE)
as.character(virus$TimeStamp[which.max(virus$interest)])
```

```
## [1] "2020-01-25T18"
```

The largest `interest` value occurred at 2020-01-25T18 which is 6:00-7:00pm on January 25, 2020. This par-

ticular hour is significant because January 25th is when the first confirmed Canadian case of the coronavirus was announced, and for many people this announcement was probably heard on the 6 o'clock news.

(b) In this question you will fit the simple linear regression model

$$y_u = \alpha + \beta x_u + r_u, \quad u \in \mathcal{P}$$

using the least squares objective function

$$\rho(\theta; \mathcal{P}) = \sum_{u \in \mathcal{P}} r_u^2$$

where $\theta = (\alpha, \beta)^T$ and $r_u = y_u - \alpha - \beta x_u$.

- i. [2 points] By taking appropriate derivatives, determine the 2×1 gradient vector $g = \nabla \rho(\theta; \mathcal{P})$. Show your work.

Solution:

$$\begin{aligned} g = \nabla \rho(\theta; \mathcal{P}) &= \begin{bmatrix} \frac{\partial \rho(\theta; \mathcal{P})}{\partial \alpha} \\ \frac{\partial \rho(\theta; \mathcal{P})}{\partial \beta} \end{bmatrix} \\ &= \begin{bmatrix} -2 \sum_{u \in \mathcal{P}} (y_u - \alpha - \beta x_u) \\ -2 \sum_{u \in \mathcal{P}} (y_u - \alpha - \beta x_u) \times x_u \end{bmatrix} \\ &= -2 \sum_{u \in \mathcal{P}} \begin{bmatrix} (y_u - \alpha - \beta x_u) \\ (y_u - \alpha - \beta x_u) \times x_u \end{bmatrix} \end{aligned}$$

- ii. [4 points] Write *factory functions* `createLSRho(x,y)` and `createLSGradient(x,y)` which take in as input only the data and which return as output the least squares objective function and the corresponding gradient function, respectively.

Solution:

```
createLSRho <- function(x,y) {
  ## Return this function
  function(theta) {
    alpha <- theta[1]
    beta <- theta[2]
    sum( (y - alpha - beta * x)^2 )
  }
}

createLSGradient <- function(x,y) {
  ## Return this function
  function(theta) {
    alpha <- theta[1]
    beta <- theta[2]
    ru = y - alpha - beta*x
    -2*c( sum(ru), sum(ru*x) )
  }
}
```

- iii. [2 point] Using the `optim` function with the `rho` and `gradient` functions created by your factory functions from part ii., find $\hat{\theta} = (\hat{\alpha}, \hat{\beta})$, the solution to

$$\underset{\theta \in \mathbb{R}^2}{\operatorname{argmin}} \rho(\theta; \mathcal{P})$$

Start the optimization at $\hat{\theta}_0 = (0, 0)$. For full points be sure to include the output from the `optim` function.

Solution:

```
rho <- createLSRrho(x = virus$h, y = virus$interest)

gradient <- createLSGradient(x = virus$h, y = virus$interest)

res_1 <- optim(par = c(0,0), fn = rho, gr = gradient)

print(res_1)

## $par
## [1] 5.0546098 0.3021034
##
## $value
## [1] 12781.38
##
## $counts
## function gradient
##      83      NA
##
## $convergence
## [1] 0
##
## $message
## NULL
```

- (c) In this question you will fit the simple linear regression model

$$y_u = \alpha + \beta x_u + r_u, \quad u \in \mathcal{P}$$

using the Huber objective function:

$$\rho(\theta; \mathcal{P}) = \sum_{u \in \mathcal{P}} \rho_k(r_u)$$

where $\theta = (\alpha, \beta)^T$, $r_u = y_u - \alpha - \beta x_u$ and

$$\rho_k(r) = \begin{cases} \frac{1}{2}r^2 & \text{for } |r| \leq k \\ k|r| - \frac{1}{2}k^2 & \text{for } |r| > k \end{cases}$$

- i. [2 points] By taking appropriate derivatives, determine the 2×1 gradient vector $g = \nabla \rho(\theta; \mathcal{P})$. Show your work.

Solution:

$$\begin{aligned}
g = \nabla \rho(\theta; \mathcal{P}) &= \begin{bmatrix} \frac{\partial \rho(\theta; \mathcal{P})}{\partial \alpha} \\ \frac{\partial \rho(\theta; \mathcal{P})}{\partial \beta} \end{bmatrix} \\
&= \begin{bmatrix} \sum_{u \in \mathcal{P}} \frac{\partial \rho_k(r_u)}{\partial \alpha} \\ \sum_{u \in \mathcal{P}} \frac{\partial \rho_k(r_u)}{\partial \beta} \end{bmatrix} \\
&= \begin{bmatrix} \sum_{u \in \mathcal{P}} \frac{\partial \rho_k(r_u)}{\partial r_u} \times \frac{\partial r_u}{\partial \alpha} \\ \sum_{u \in \mathcal{P}} \frac{\partial \rho_k(r_u)}{\partial r_u} \times \frac{\partial r_u}{\partial \beta} \end{bmatrix} \\
&= \begin{bmatrix} -\sum_{u \in \mathcal{P}} \frac{\partial \rho_k(r_u)}{\partial r_u} \\ -\sum_{u \in \mathcal{P}} \frac{\partial \rho_k(r_u)}{\partial r_u} \times x_u \end{bmatrix} \\
&= -\sum_{u \in \mathcal{P}} \begin{bmatrix} \frac{\partial \rho_k(r_u)}{\partial r_u} \\ \frac{\partial \rho_k(r_u)}{\partial r_u} \times x_u \end{bmatrix}
\end{aligned}$$

where

$$\frac{\partial \rho_k(r)}{\partial r} = \begin{cases} r & \text{for } |r| \leq k \\ \text{sign}(r) \times k & \text{for } |r| > k \end{cases}$$

- ii. [4 points] Write *factory functions* `createHuberRho(x,y,k)` and `createHuberGradient(x,y,k)` which take in as input only the data and the threshold k , and which return as output the Huber objective function and the corresponding gradient function, respectively. Note that you may use the `huber.fn` and `huber.fn.prime` functions from class as necessary.

Solution:

```

createHuberRho <- function(x,y,k) {
  ## Return this function
  function(theta) {
    alpha <- theta[1]
    beta <- theta[2]
    sum( huber.fn(y - alpha - beta * x, k = k ) )
  }
}

createHuberGradient <- function(x,y,k) {
  ## Return this function
  function(theta) {
    alpha <- theta[1]
    beta <- theta[2]
    ru = y - alpha - beta*x
    rhok = huber.fn.prime(ru, k=k)
    -1*c( sum(rhok*1), sum(rhok*x) )
  }
}

```

- iii. [2 points] Using the `optim` function with the `rho` and `gradient` functions created by your factory functions from part ii., find $\hat{\theta} = (\hat{\alpha}, \hat{\beta})$, the solution to

$$\underset{\theta \in \mathbb{R}^2}{\operatorname{argmin}} \rho(\theta; \mathcal{P})$$

while letting $k = 6.268$ and $\hat{\theta}_0 = (0, 0)$. For full points be sure to include the output from the `optim` function.

Solution:

```
rho <- createHuberRho(x = virus$h, y = virus$interest, k = 6.268)

gradient <- createHuberGradient(x = virus$h, y = virus$interest, k = 6.268)

res_r <- optim(par = c(0,0), fn = rho, gr = gradient)

print(res_r)

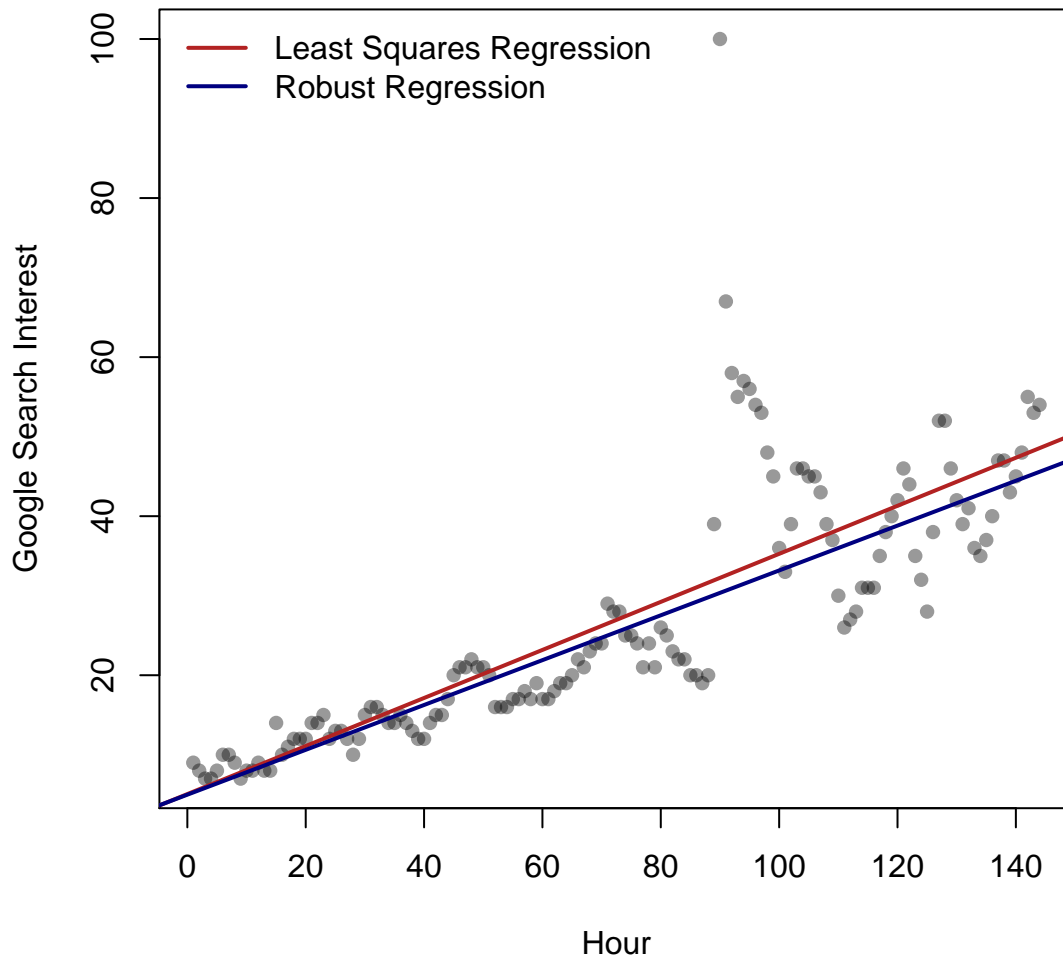
## $par
## [1] 4.9796283 0.2819384
##
## $value
## [1] 3002.395
##
## $counts
## function gradient
##      93      NA
##
## $convergence
## [1] 0
##
## $message
## NULL
```

- (d) [5 points] Recreate the scatter plot from above and overlay the least squares line and the robust regression line, distinguished by a legend.

Solution:

```
plot(x = virus$h, y = virus$interest, pch = 16, col = adjustcolor("black", 0.4),
     ylab = "Google Search Interest", xlab = "Hour",
     main = "Popularity of the Search Term \"Coronavirus\" over Time")
abline(res_l$par, col = "firebrick", lwd = 2)
abline(res_r$par, col = "navyblue", lwd = 2)
legend("topleft", col = c("firebrick", "navyblue"), lty = 1, lwd = 2, bty = "n",
     legend = c("Least Squares Regression", "Robust Regression"))
```

Popularity of the Search Term "Coronavirus" over Time



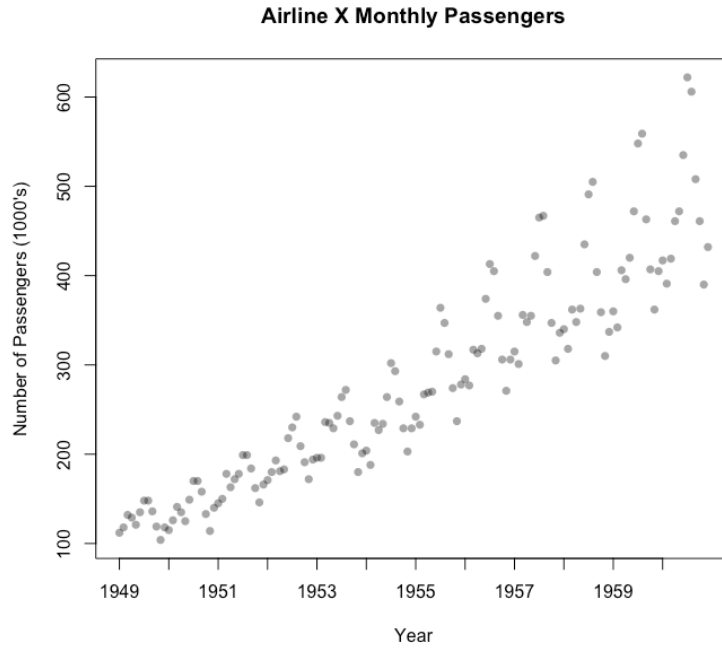
- (e) [2 points] Does the outlier identified in part (a) appear to have a large influence on the least squares line? State YES or NO and provide a brief justification.

Solution: NO. Although the least squares and robust regression lines are not exactly the same, they are not dramatically different suggesting that 2020-01-25T18 did not have a large influence.

QUESTION 3: Modeling Airline Ticket Sales [27 points]

The `AirPassengers` dataset is a classic dataset used in times series and forecasting. For Airline X (the actual airline is kept confidential due to privacy concerns) this dataset contains $N = 144$ observations corresponding to the monthly totals of passengers (in thousands) from January 1949 to December 1960. The data is available in the `airpassengers.json` file and summarized and visualized below.

Variate	Description
<code>month</code>	A numeric variate indicating month. Note that <code>month=1</code> corresponds to January 1949, <code>month=2</code> corresponds to February 1949 and so on.
<code>passengers</code>	A numeric variate indicating the number of passengers (in 1000s) on Airline X flights in a given month.



(a) In this question you will fit the simple linear regression model

$$y_u = \alpha + \beta x_u + r_u, \quad u \in \mathcal{P}$$

via the Newton-Raphson Method in conjunction with the least squares objective function

$$\rho(\theta; \mathcal{P}) = \sum_{u \in \mathcal{P}} r_u^2$$

where $\theta = (\alpha, \beta)^T$ and $r_u = y_u - \alpha - \beta x_u$.

- i. [4 points] Determine the vector $\psi(\theta; \mathcal{P})$ and matrix $\psi'(\theta; \mathcal{P})$ required to apply the Newton-Raphson method. Show your work.

Solution:

$$\begin{aligned} \psi(\theta; \mathcal{P}) &= \nabla \rho(\theta; \mathcal{P}) = \begin{bmatrix} \frac{\partial \rho(\theta; \mathcal{P})}{\partial \alpha} \\ \frac{\partial \rho(\theta; \mathcal{P})}{\partial \beta} \end{bmatrix} = -2 \begin{bmatrix} \sum_{u \in \mathcal{P}} (y_u - \alpha - \beta x_u) \\ \sum_{u \in \mathcal{P}} (y_u - \alpha - \beta x_u) \times x_u \end{bmatrix} = \begin{bmatrix} \psi_1 \\ \psi_2 \end{bmatrix} \\ \psi'(\theta; \mathcal{P}) &= \begin{bmatrix} \frac{\partial \psi_1}{\partial \alpha} & \frac{\partial \psi_1}{\partial \beta} \\ \frac{\partial \psi_2}{\partial \alpha} & \frac{\partial \psi_2}{\partial \beta} \end{bmatrix} = -2 \begin{bmatrix} \sum_{u \in \mathcal{P}} (-1) & \sum_{u \in \mathcal{P}} (-x_u) \\ \sum_{u \in \mathcal{P}} (-x_u) & \sum_{u \in \mathcal{P}} (-x_u^2) \end{bmatrix} = 2 \begin{bmatrix} N & \sum_{u \in \mathcal{P}} x_u \\ \sum_{u \in \mathcal{P}} x_u & \sum_{u \in \mathcal{P}} x_u^2 \end{bmatrix} \end{aligned}$$

- ii. [4 points] Write *factory functions* `createLSpsi(x,y)` and `createLSpsiPrime(x,y)` which take in as input only the data and which return as output the `psi` and `psiPrime` functions associated with the least squares objective function (for the linear model).

Solution:

```
createLSpsi <- function(x,y){
  ## return this function:
  function(theta){
    alpha <- theta[1]
    beta <- theta[2]
```

```

    ru = y - alpha - beta*x
    -2*matrix(c( sum(ru), sum(ru*x) ), nrow = 2)
  }
}

createLSpsiPrime <- function(x,y){
  ## return this function:
  function(theta){
    alpha <- theta[1]
    beta <- theta[2]
    2*matrix(c(length(x), sum(x), sum(x), sum(x^2)), nrow = 2, ncol = 2)
  }
}

```

- iii. [2 points] Using the `NewtonRaphson` function (from class) together with the `testConvergence` function (from class) as well as `psi` and `psiPrime` functions created by your factory functions from part ii., find $\hat{\theta} = (\hat{\alpha}, \hat{\beta})$, the solution to $\psi(\theta; \mathcal{P}) = 0$. Start the optimization at $\hat{\theta}_0 = (0, 0)$. For full points be sure to include the output from the `NewtonRaphson` function.

Solution:

```

library(rjson)
flights <- as.data.frame(fromJSON(file = "/Users/nstevens/Dropbox/Teaching/STAT_341/Assignments/Assignm

psi <- createLSpsi(x = flights$month, y = flights$passengers)
psiPrime <- createLSpsiPrime(x = flights$month, y = flights$passengers)
res_l_NR <- NewtonRaphson(theta= c(0,0), psiFn = psi, psiPrimeFn = psiPrime)
print(res_l_NR)

## $theta
##           [,1]
## [1,] 87.652778
## [2,]  2.657184
##
## $converged
## [1] TRUE
##
## $iteration
## [1] 2
##
## $fnValue
##           [,1]
## [1,] -1.719513e-12
## [2,] -2.004370e-10

```

- (b) In this question you will fit the linear regression model

$$y_u = \alpha + \beta x_u + \gamma x_u^2 + r_u, \quad u \in \mathcal{P}$$

via the Newton-Raphson Method in conjunction with the least squares objective function

$$\rho(\theta; \mathcal{P}) = \sum_{u \in \mathcal{P}} r_u^2$$

where $\theta = (\alpha, \beta, \gamma)^T$ and $r_u = y_u - \alpha - \beta x_u - \gamma x_u^2$.

- i. [4 points] Determine the vector $\psi(\theta; \mathcal{P})$ and matrix $\psi'(\theta; \mathcal{P})$ required to apply the Newton-Raphson method. Show your work.

Solution:

$$\psi(\theta; \mathcal{P}) = \nabla \rho(\theta; \mathcal{P}) = \begin{bmatrix} \frac{\partial \rho(\theta; \mathcal{P})}{\partial \alpha} \\ \frac{\partial \rho(\theta; \mathcal{P})}{\partial \beta} \\ \frac{\partial \rho(\theta; \mathcal{P})}{\partial \gamma} \end{bmatrix} = -2 \begin{bmatrix} \sum_{u \in \mathcal{P}} (y_u - \alpha - \beta x_u - \gamma x_u^2) \\ \sum_{u \in \mathcal{P}} (y_u - \alpha - \beta x_u - \gamma x_u^2) \times x_u \\ \sum_{u \in \mathcal{P}} (y_u - \alpha - \beta x_u - \gamma x_u^2) \times x_u^2 \end{bmatrix} = \begin{bmatrix} \psi_1 \\ \psi_2 \\ \psi_3 \end{bmatrix}$$

$$\psi'(\theta; \mathcal{P}) = \begin{bmatrix} \frac{\partial \psi_1}{\partial \alpha} & \frac{\partial \psi_1}{\partial \beta} & \frac{\partial \psi_1}{\partial \gamma} \\ \frac{\partial \psi_2}{\partial \alpha} & \frac{\partial \psi_2}{\partial \beta} & \frac{\partial \psi_2}{\partial \gamma} \\ \frac{\partial \psi_3}{\partial \alpha} & \frac{\partial \psi_3}{\partial \beta} & \frac{\partial \psi_3}{\partial \gamma} \end{bmatrix} = 2 \begin{bmatrix} N & \sum_{u \in \mathcal{P}} x_u & \sum_{u \in \mathcal{P}} x_u^2 \\ \sum_{u \in \mathcal{P}} x_u & \sum_{u \in \mathcal{P}} x_u^2 & \sum_{u \in \mathcal{P}} x_u^3 \\ \sum_{u \in \mathcal{P}} x_u^2 & \sum_{u \in \mathcal{P}} x_u^3 & \sum_{u \in \mathcal{P}} x_u^4 \end{bmatrix}$$

- ii. [4 points] Write *factory functions* `createLSQpsi(x,y)` and `createLSQpsiPrimew(x,y)` which take in as input only the data and which return as output the `psi` and `psiPrime` functions associated with the least squares objective function (for the quadratic model).

Solution:

```
```r
createLSQpsi <- function(x,y) {
 ## Return this function
 function(theta) {
 alpha <- theta[1]
 beta <- theta[2]
 gamma <- theta[3]
 ru = y - alpha - beta*x - gamma*x^2
 -2*matrix(c(sum(ru), sum(ru*x), sum(ru*x^2)), nrow = 3)
 }
}

createLSQpsiPrime <- function(x,y) {
 ## Return this function
 function(theta) {
 alpha <- theta[1]
 beta <- theta[2]
 gamma <- theta[3]
 2*matrix(c(length(x), sum(x), sum(x^2),
 sum(x), sum(x^2), sum(x^3),
 sum(x^2), sum(x^3), sum(x^4)), nrow = 3, ncol = 3)
 }
}
```
```

- iii. [2 points] Using the `NewtonRaphson` function (from class) together with the `testConvergence` function (from class) as well as `psi` and `psiPrime` functions created by your factory functions from part ii., find $\hat{\theta} = (\hat{\alpha}, \hat{\beta}, \hat{\gamma})$, the solution to $\psi(\theta; \mathcal{P}) = 0$. Start the optimization at $\hat{\theta}_0 = (0, 0, 0)$. For full points be sure to include the output from the `NewtonRaphson` function.

Solution:

```
psi_q <- createLSQpsi(x = flights$month, y = flights$passengers)
psiPrime_q <- createLSQpsiPrime(x = flights$month, y = flights$passengers)
res_q_NR <- NewtonRaphson(theta = c(0,0,0), psiFn = psi_q, psiPrimeFn = psiPrime_q)
print(res_q_NR)
```

```
## $theta
##           [,1]
## [1,] 1.123800e+02
## [2,] 1.640995e+00
## [3,] 7.008198e-03
##
## $converged
## [1] TRUE
##
## $iteration
## [1] 2
##
## $fnValue
##           [,1]
## [1,] 1.448230e-12
## [2,] 9.128698e-11
## [3,] 7.062226e-09
```

- (c) [5 points] Recreate the scatter plot from above and overlay the least squares line and the least squares quadratic curve, distinguished by a legend.

Solution:

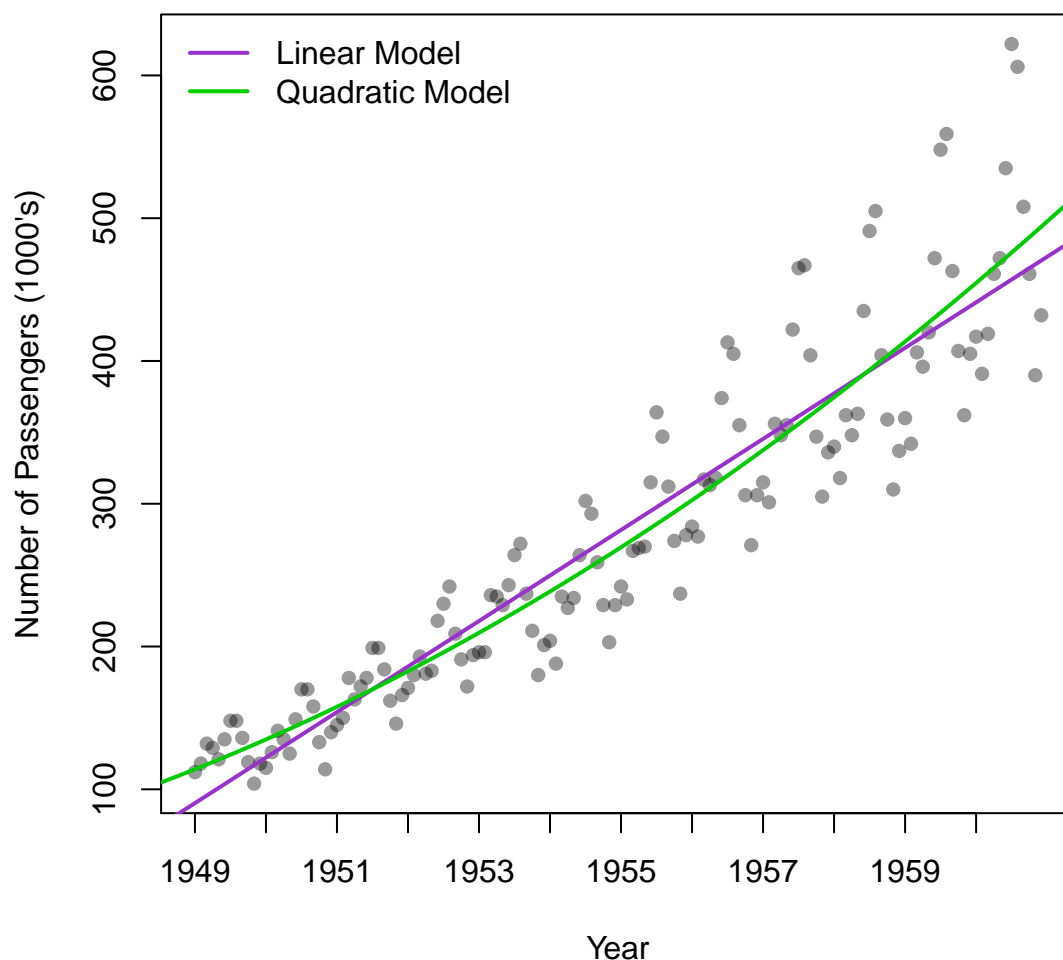
```
plot(x = flights$month, y = flights$passengers, pch = 16, col = adjustcolor("black", 0.4),
     xlab = "Year", ylab = "Number of Passengers (1000's)",
     main = "Airline X Monthly Passengers", xaxt = "n")
axis(side = 1, at = (12*0:11)+1, labels = 1949:1960)

abline(a = res_l_NR$theta[1], b = res_l_NR$theta[2], col = "darkorchid", lwd = 2)

quad <- res_q_NR$theta[1] +
  res_q_NR$theta[2]*seq(-10, 150, 0.1) +
  res_q_NR$theta[3]*seq(-10, 150, 0.1)^2
lines(x = seq(-10, 150, 0.1), y = quad, col = "green3", lwd = 2)

legend("topleft", col = c("darkorchid", "green3"), lty = 1, lwd = 2, bty = "n",
      legend = c("Linear Model", "Quadratic Model"))
```

Airline X Monthly Passengers



- (d) [2 points] Based on the plot in part (c), which model best represents the relationship in the population – the line or the quadratic curve? Briefly justify your response.

Solution: The quadratic curve appears to fit the data better. The performance of the curves are not very different in the middle of the data, but the quadratic curve does a better job summarizing the pattern in the tails of the data (the early and later years).