

STAT 341: Tutorial 4 – Practice with Implicit Attributes

Friday January 31, 2020

Part I: Not all Estimating Equations Arise from Derivatives

With this example I want to make clear that although many of the estimating equations we encounter arise as derivatives from some objective function, the *don't have to*.

Consider the implicitly defined attribute $\theta = (\mu, \sigma)$ where μ and σ are respectively measures of center and spread in the population $\mathcal{P} = \{y_1, y_2, \dots, y_N\}$. Find $\hat{\theta} = (\hat{\alpha}, \hat{\beta})$, the solution to the following system of equations

$$\psi(\theta; \mathcal{P}) = \mathbf{0}$$

given by

$$\begin{bmatrix} \left(\frac{1}{N} \sum_{u \in \mathcal{P}} y_u \right) - \mu \\ \left(\frac{1}{N} \sum_{u \in \mathcal{P}} y_u^2 \right) - \mu^2 - \sigma^2 \end{bmatrix} = \begin{bmatrix} 0 \\ 0 \end{bmatrix}$$

Note that this system of equations arose by equating sample moments with population moments. This method of estimation is referred to as the Method of Moments.

Part II: Least Absolute Deviations Regression

In class we have talked a lot about estimating $\theta = (\alpha, \beta)$ the intercept and slope associated with the simple linear regression

$$y_u = \alpha + \beta(x_u - \bar{x}) + r_u$$

And we have done this in a variety of different ways by altering the objective function. One such possible objective function is the *absolute error loss* function which gives rise to **least absolute deviations (LAD) regression**:

$$\hat{\theta} = \underset{\theta \in \Theta}{\operatorname{argmin}} \sum_{u \in \mathcal{P}} |y_u - \alpha - \beta(x_u - \bar{x})|$$

In this part we will use gradient descent to fit this model to the **Animals** data. But first we need to define **rho** and **gradient** functions in R which requires that we first compute the gradients by hand.

(a) Let

$$\rho(\theta; \mathcal{P}) = \sum_{u \in \mathcal{P}} |y_u - \alpha - \beta(x_u - \bar{x})|$$

Calculate the gradient vector $\mathbf{g} = \nabla \rho(\theta; \mathcal{P})$.

- (b) Write *factory functions* `createLADRho(x,y)` and `createLADGradient(x,y)` which take in as input only the data and which return as output the least absolute deviations objective function and the corresponding gradient function, respectively

```
createLADRho <- function(x, y) {  
  ## local variable  
  xbar <- mean(x)  
  ## Return this function  
  function(theta) {  
    alpha <- theta[1]  
    beta <- theta[2]  
    sum(abs(y - alpha - beta * (x - xbar)))  
  }  
}  
  
createLADGradient <- function(x, y) {  
  ## local variables  
  xbar <- mean(x)  
  function(theta) {  
    alpha <- theta[1]  
    beta <- theta[2]  
    ru = y - alpha - beta * (x - xbar)  
    -1 * c(sum(sign(ru)), sum(sign(ru) * (x - xbar)))  
  }  
}
```

- (c) Using the `gradientDescent` function (from class) together with the `gridLineSearch` and `testConvergence` functions (from class) as well as `rho` and `gradient` functions created by your factory functions from part (b), find the LAD estimates $\hat{\theta} = (\hat{\alpha}, \hat{\beta})$ for the `Animals` data.

```
library(robustbase)
rho <- createLADrho(x = log(Animals2$body), y = log(Animals2$brain))

g <- createLADgradient(x = log(Animals2$body), y = log(Animals2$brain))

res.manual <- gradientDescent(theta = c(0, 0), rhoFn = rho, gradientFn = g,
  lineSearchFn = gridLineSearch, testConvergenceFn = testConvergence,
  maxIterations = 5000, tolerance = 1e-20, relative = TRUE)

print(res.manual)

## $theta
## [1] 3.3589387 0.7368431
##
## $converged
## [1] TRUE
##
## $iteration
## [1] 26
##
## $fnValue
## [1] 48.01975
```

- (d) In class we performed LAD regression using the `lad` function from the `L1pack` package. Let's confirm that what we found above agrees with the output of this other function.

```
library(L1pack)
res.L1pack <- lad(log(Animals2$brain) ~ I(log(Animals2$body) - mean(log(Animals2$body))))
print(res.L1pack)

## Call:
## lad(formula = log(Animals2$brain) ~ I(log(Animals2$body) - mean(log(Animals2$body))))
## Converged in 9 iterations
##
## Coefficients:
##                                (Intercept)
##                                3.3541
## I(log(Animals2$body) - mean(log(Animals2$body)))
##                                0.7373
##
## Degrees of freedom: 65 total; 63 residual
## Scale estimate: 1.04454
```


- (e) Construct a scatter plot of $\log(\text{brain weight})$ versus $\log(\text{body weight})$ for the `Animals` data and plot both lines of best fit – the one we determined manually and the one calculated using `lad` – to see if the difference is material. Use a legend to distinguish among the lines.

```
plot(x = log(Animals2$body), y = log(Animals2$brain), main = "", xlab = "log(body weight)",
     ylab = "log(brain weight)", pch = 16, col = adjustcolor("black", alpha.f = 0.3))

abline(a = res.manual$theta[1] - res.manual$theta[2] * mean(log(Animals2$body)),
       b = res.manual$theta[2], col = "firebrick", lwd = 2, lty = 1)

abline(a = res.L1pack$coef[1] - res.L1pack$coef[2] * mean(log(Animals2$body)),
       b = res.L1pack$coef[2], col = "navyblue", lwd = 2, lty = 2)

legend("bottomright", legend = c("LAD (manual)", "LAD (L1pack)"), col = c("firebrick",
                                "navyblue"), lty = 1:2, lwd = 2)
```



- (f) Can we apply the Newton-Raphson Method to this problem? In other words, can we determine the LAD estimate $\hat{\boldsymbol{\theta}} = (\hat{\alpha}, \hat{\beta})$ via the Newton-Raphson Method?