

# STAT 341: Tutorial 8 – Practice with Bootstrap & Prediction

Friday March 27, 2020


## Part I: Bootstrap Confidence Intervals

Here we will deal with the Billboard Top 30 data again, and in particular the popularity score of the  $N = 300$  songs.

(a) Load the data and plot a histogram of the popularity score.

```
spot <- read.csv("/Users/nstevens/Dropbox/Teaching/STAT_341/Assignments/Assignment1/spotify.csv")
hist(spot$popularity, xlab = "Popularity Score", main = "Billboard Top 30, 2010-2019",
     col = adjustcolor("grey", 0.4))
```

**Billboard Top 30, 2010–2019**




(b) Take a sample of size  $n = 50$  from this population and plot the popularity scores from this sample.

```
# Take the sample
set.seed(341)
N <- dim(spot)[1]
n <- 50
samp.indx <- sample(N, n)
S <- spot[samp.indx, "popularity"]
# Make the plot
```

```
hist(S, xlab = "Popularity Score", main = "Billboard Top 30, Sample (n=50)",
     col = adjustcolor("grey", 0.4))
```

### Billboard Top 30, Sample (n=50)



- (c) By resampling  $\mathcal{S}$  with replacement, construct  $B = 5000$  bootstrap samples  $S_1^*, S_2^*, \dots, S_{5000}^*$  and construct histograms of the average and interquartile range calculated on each of those samples. Include vertical lines that indicate the average and IQR in the population.

```
# Get the bootstrap samples
B <- 5000
Sstar <- sapply(1:B, FUN = function(b) {
  sample(S, n, replace = TRUE)
})

# Calculate the average on all of these bootstrap samples
avg_star <- apply(X = Sstar, MARGIN = 2, FUN = mean)


# Calculate the IQR on all of these bootstrap samples
iqr_star <- apply(X = Sstar, MARGIN = 2, FUN = IQR)

# Construct histograms for each of these
par(mfrow = c(1, 2))

hist(avg_star, col = adjustcolor("darkgreen", 0.5), xlab = "Average Popularity Score",
     main = "5000 Bootstrap Replicates")
abline(v = mean(spot$popularity), col = "purple", lwd = 2)

hist(iqr_star, col = adjustcolor("darkgreen", 0.5), xlab = "IQR of Popularity Scores",
```

```
main = "5000 Bootstrap Replicates")
abline(v = IQR(spot$popularity), col = "purple", lwd = 2)
```



- (d) Calculate naive normal theory, quantile method, and bootstrap-t 90% confidence intervals for the population average. For the bootstrap-t interval use  $B = 5000$  and  $D = 100$ .

```
bootstrap_t_interval <- function(S, a, confidence, B, D) {
  ## Inputs: S = an n element array containing the variate values in the
  ## sample a = a scalar-valued function that calculates the attribute a()
  ## of interest confidence = a value in (0,1) indicating the confidence
  ## level B = a numeric value representing the outer bootstrap count of
  ## replicates (used to calculate the lower and upper limits) D = a
  ## numeric value representing the inner bootstrap count of replicates
  ## (used to estimate the standard deviation of the sample attribute for
  ## each (outer) bootstrap sample)

  Pstar <- S
  aPstar <- a(Pstar)
  sampleSize <- length(S)
  ## get (outer) bootstrap values
  bVals <- sapply(1:B, FUN = function(b) {
    Sstar <- sample(Pstar, sampleSize, replace = TRUE)
    aSstar <- a(Sstar)
    ## get (inner) bootstrap values to estimate the SD
    Pstarstar <- Sstar
    SD_aSstar <- sd(sapply(1:D, FUN = function(d) {
      Sstarstar <- sample(Pstarstar, sampleSize, replace = TRUE)
      ## return the attribute value
      a(Sstarstar)
    }))
    z <- (aSstar - aPstar)/SD_aSstar
    ## Return the two values
    c(aSstar = aSstar, z = z)
  })
}
```

```

SDhat <- sd(bVals["aSstar", ])
zVals <- bVals["z", ]
## Now use these zVals to get the lower and upper c values.
cValues <- quantile(zVals, probs = c((1 - confidence)/2, (confidence +
1)/2), na.rm = TRUE)
cLower <- min(cValues)
cUpper <- max(cValues)
interval <- c(lower = aPstar - cUpper * SDhat, middle = aPstar, upper = aPstar -
cLower * SDhat)
interval

```

$$[a(s) - c \times \hat{SD}_*[\tilde{a}(s)], a(s) + c \times \hat{SD}_*[\tilde{a}(s)]]$$

$c$  is determined from  $Z \sim N(0,1)$

*Naive normal theory*  
 $\text{mean}(S) + \text{qnorm}(0.95) * c(-1, 1) * \text{sd}(\text{avg\_star})$

```
## [1] 75.50064 78.01936
```

*Quantile method*

```
c(quantile(avg_star, 0.05), quantile(avg_star, 0.95))
```

```
## 5% 95%
```

```
## 75.54 78.06
```

*Bootstrap-t*

```
bootstrap_t_interval(S = S, a = mean, confidence = 0.9, B = 5000, D = 100)
```

```
## lower middle upper
```

```
## 75.61863 76.76000 78.15993
```

(e) Calculate naive normal theory, quantile method, and bootstrap-t 90% confidence intervals for the population IQR. For the bootstrap-t interval use  $B = 5000$  and  $D = 100$ .

*Naive normal theory*

```
IQR(S) + qnorm(0.95) * c(-1, 1) * sd(iqr_star)
```

```
## [1] 4.16858 7.83142
```

*Quantile method*

```
c(quantile(iqr_star, 0.05), quantile(iqr_star, 0.95))
```

```
## 5% 95%
```

```
## 3.75 7.50
```

*Bootstrap-t*

```
bootstrap_t_interval(S = S, a = IQR, confidence = 0.9, B = 5000, D = 100)
```

```
## lower middle upper
```

```
## 4.702215 6.000000 8.485624
```


## Part II: Predictions and APSE Decomposition

Here we will use some simulated data as our population, from which we will draw a sample, fit various models, and evaluate them on a test set. The data can be found in the `fakedata.csv` file and the relevant variates are  $x$  and  $y$ .

$$\text{APSE}(\mathcal{P}, \tilde{\mu}_s(x)) = \underbrace{\text{Ave}_x(\text{Var}(y(x)) + \text{Var}[\tilde{\mu}_s(x)] + \text{Bias}[\tilde{\mu}_s(x)]^2)}_{\text{irreducible error}}$$

(a) Load and visualize the data.

```
fake <- read.csv("/Users/nstevens/Dropbox/Teaching/STAT_341/Tutorials/Tutorial 8/fakedata.csv",
  header = TRUE)
plot(fake, pch = 16, col = adjustcolor("darkgreen", 0.5))
```



(b) Recreate the plot above, but this time overlay linear, quadratic and quartic polynomials, fitted using all of the data. Make sure the curves have different colours and are distinguished with a legend. **Note:** the getmuhat function will be useful.

```
# Define useful function
getmuhat <- function(sampleXY, complexity = 1){
  formula <- paste0("y ~ ",
    if (complexity==0) {
      "1"
    } else {
      paste0("poly(x, ", complexity, ", raw = FALSE)")
    }
    #paste0("bs(x, ", complexity, ")")
  )

  fit <- lm(as.formula(formula), data = sampleXY)
  tx = sampleXY$x
  ty = fit$fitted.values

  range.X = range(tx)
  val.rY = c( mean(ty[tx == range.X[1]]),
    mean(ty[tx == range.X[2]]) )

  ## From this we construct the predictor function
  muhat <- function(x){
    if ("x" %in% names(x)) {
```

```

    ## x is a dataframe containing the variate named
    ## by xvarname
    newdata <- x
  } else
    ## x is a vector of values that needs to be a data.frame
    { newdata <- data.frame(x = x) }
  ## The prediction
  ##
  val = predict(fit, newdata = newdata)
  val[newdata$x < range.X[1]] = val.rY[1]
  val[newdata$x > range.X[2]] = val.rY[2]
  val
}
## muhat is the function that we need to calculate values
## at any x, so we return this function from getmuhat
muhat
}

```

```

# Make the plot
plot(fake, pch = 16, col = adjustcolor("darkgreen", 0.5))


# Fit a linear polynomial and add it to the plot
muhat1 <- getmuhat(fake, 1)
curve(muhat1, from = 0, to = 1, add = TRUE, col = "black", lwd = 3)

# Fit a quadratic polynomial and add it to the plot
muhat2 <- getmuhat(fake, 2)
curve(muhat2, from = 0, to = 1, add = TRUE, col = "blue", lwd = 3)

# Fit a quartic polynomial and add it to the plot
muhat4 <- getmuhat(fake, 4)
curve(muhat4, from = 0, to = 1, add = TRUE, col = "darkred", lwd = 3)

# Add a legend
legend("topright", legend = c("deg = 1", "deg = 2", "deg = 4"), col = c("black",
  "blue", "darkred"), lwd = 3, cex = 0.8, bty = "n")

```



- (c) Generate  $N_S = 100$  samples of size  $n = 40$  and fit polynomials of degree 1, 2 and 4 to every sample.  
**Note:** the `getSampleComp` and `getXYSample` functions will be useful.

```
# Define useful functions
getSampleComp <- function(pop, size, replace = FALSE) {
  N <- nrow(as.data.frame(pop))
  samp <- rep(FALSE, N)
  samp[sample(1:N, size, replace = replace)] <- TRUE
  samp
}

getXYSample <- function(xvarname, yvarname, samp, pop) {
  sampData <- pop[samp, c(xvarname, yvarname)]
  names(sampData) <- c("x", "y")
  sampData
}
```

```
N_S <- 100
n <- 40
set.seed(341)
```

```
# Use getSampleComp via lapply to determine which units to include in
# the 100 samples. samps is a list with 100 elements, each of which is
# a 250-element array, 40 of which are TRUE and 210 of which are FALSE.
# This tells us which units to include in each of the 100 samples
samps <- lapply(1:N_S, FUN = function(i) {
  getSampleComp(fake, n)
})
```

```
# getXYsample is applied to fake and each element of samps. The code
```

```
# below returns Ssamples which is another 100-element list, but each of
# these elements is a data frame with columns x and y and containing
# just the sample data.
```

```
Ssamples <- lapply(samps, FUN = function(Si) {
  getXYSample("x", "y", Si, fake)
})
```

```
# Tsamples is the complement of Ssamples; it is a 100-element list,
# where each of the elements is a data frame with columns x and y and
# containing just the test data.
```

```
Tsamples <- lapply(samps, FUN = function(Si) {
  getXYSample("x", "y", !Si, fake)
})
```

```
# This code applies getmuhat to each element of Ssamples and thereby
# fits a polynomial of the given complexity to every sample and saves
# those results in a list
```

```
muhats1 <- lapply(Ssamples, getmuhat, complexity = 1)
muhats2 <- lapply(Ssamples, getmuhat, complexity = 2)
muhats4 <- lapply(Ssamples, getmuhat, complexity = 4)
```

- (d) Using `par(mfrow=c(1,3))` plot the data and overlay all the fitted polynomials from part (c) with degree 1 depicted in the left plot, degree 2 depicted on the middle plot, and degree 4 in the right plot. Use colours that are consistent with part (b).

```
par(mfrow = c(1, 3))
```

```
# Plot the data
```

```
plot(fake, pch = 16, col = adjustcolor("darkgreen", 0.5), main = bquote("Degree 1" ~
  hat(mu) * "'s"))
```

```
# Add the degree 1 curves
```

```
for (i in 1:N_S) {
  curveFn <- muhats1[[i]]
  curve(curveFn, from = 0, to = 1, add = TRUE, col = adjustcolor("black",
    0.25))
}
```

```
# Plot the data
```

```
plot(fake, pch = 16, col = adjustcolor("darkgreen", 0.5), main = bquote("Degree 2" ~
  hat(mu) * "'s"))
```

```
# Add the degree 2 curves
```

```
for (i in 1:N_S) {
  curveFn <- muhats2[[i]]
  curve(curveFn, from = 0, to = 1, add = TRUE, col = adjustcolor("blue",
    0.25))
}
```

```
# Plot the data
```

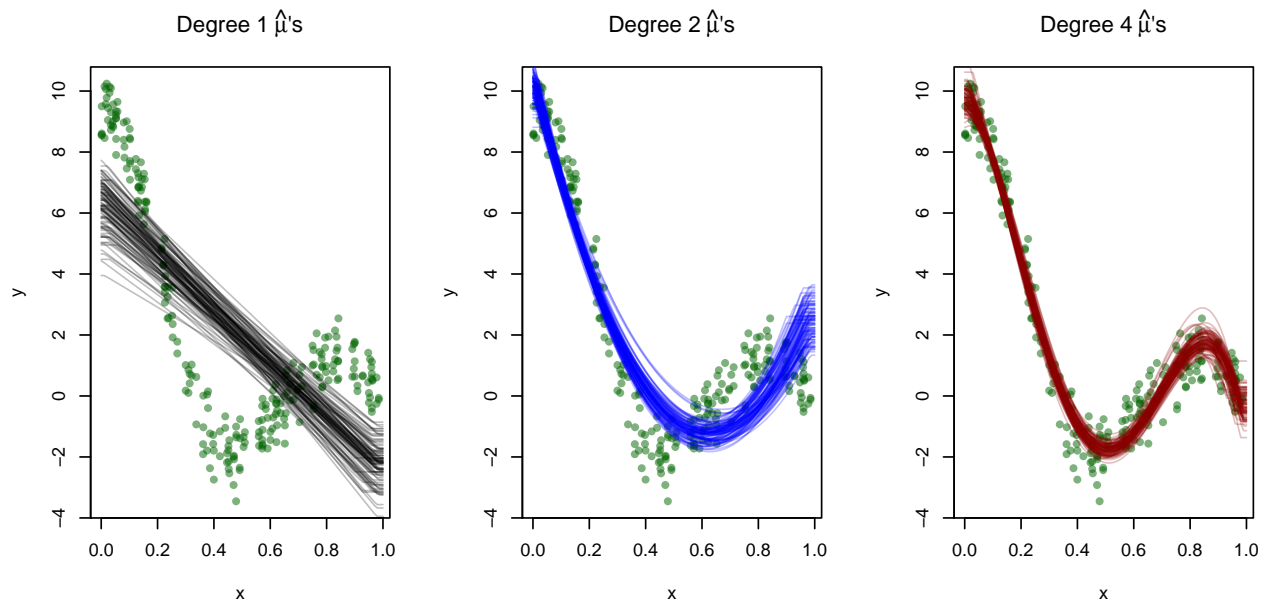
```
plot(fake, pch = 16, col = adjustcolor("darkgreen", 0.5), main = bquote("Degree 4" ~
  hat(mu) * "'s"))
```

```
# Add the degree 4 curves
```

```
for (i in 1:N_S) {
  curveFn <- muhats4[[i]]
  curve(curveFn, from = 0, to = 1, add = TRUE, col = adjustcolor("darkred",
```



```
0.25))
}
```



- (e) Using the  $N_S = 100$  samples of size  $n = 40$  generated in part (c), calculate the APSE (and each of its components) for degrees in 0:15. Summarize the results in tabular and graphical format. **Note:** the apse\_all, getmubar and getmuFun functions will be useful.

```
# Define useful functions
apse_all <- function(Ssamples, Tsamples, complexity, mu) {
  ## average over the samples S
  N_S <- length(Ssamples)
  muhats <- lapply(Ssamples, FUN = function(sample) getmuhat(sample,
    complexity))
  ## get the average of these, mubar
  mubar <- getmubar(muhats)

  rowMeans(sapply(1:N_S, FUN = function(j) {
    T_j <- Tsamples[[j]]
    muhat <- muhats[[j]]
    ## Take care of any NAs
    T_j <- na.omit(T_j)
    y <- T_j$y
    x <- T_j$x
    mu_x <- mu(x)
    muhat_x <- muhat(x)
    mubar_x <- mubar(x)

    ## apse average over (x_i, y_i) in a single sample T_j the squares (y -
    ## muhat(x))^2
    apse <- (y - muhat_x)

    ## bias2: average over (x_i, y_i) in a single sample T_j the squares (y -
    ## muhat(x))^2
    bias2 <- (mubar_x - mu_x)
```

```

    ## var_mutilde average over (x_i,y_i) in a single sample T_j the squares
    ## (y - muhat(x))^2
    var_mutilde <- (muhat_x - mubar_x)

    ## var_y : average over (x_i,y_i) in a single sample T_j the squares (y
    ## - muhat(x))^2
    var_y <- (y - mu_x)

    ## Put them together and square them
    squares <- rbind(apse, var_mutilde, bias2, var_y)^2

    ## return means
    rowMeans(squares)
  )))
}

getmubar <- function(muhats) {
  function(x) {
    Ans <- sapply(muhats, FUN = function(muhat) {
      muhat(x)
    })
    apply(Ans, MARGIN = 1, FUN = mean)
  }
}

getmuFun <- function(pop, xvarname, yvarname) {
  pop = na.omit(pop[, c(xvarname, yvarname)])

  # rule = 2 means return the nearest y-value when extrapolating, same as
  # above. ties = mean means that repeated x-values have their y-values
  # averaged, as above.
  muFun = approxfun(pop[, xvarname], pop[, yvarname], rule = 2, ties = mean)
  return(muFun)
}

muhat <- getmuFun(fake, "x", "y")
degrees <- 0:15
apse_vals <- sapply(degrees, FUN = function(complexity) {
  apse_all(Ssamples, Tsamples, complexity = complexity, mu = muhat)
})

# Print out the results in a table
t(rbind(degrees, apse = round(apse_vals, 5)))

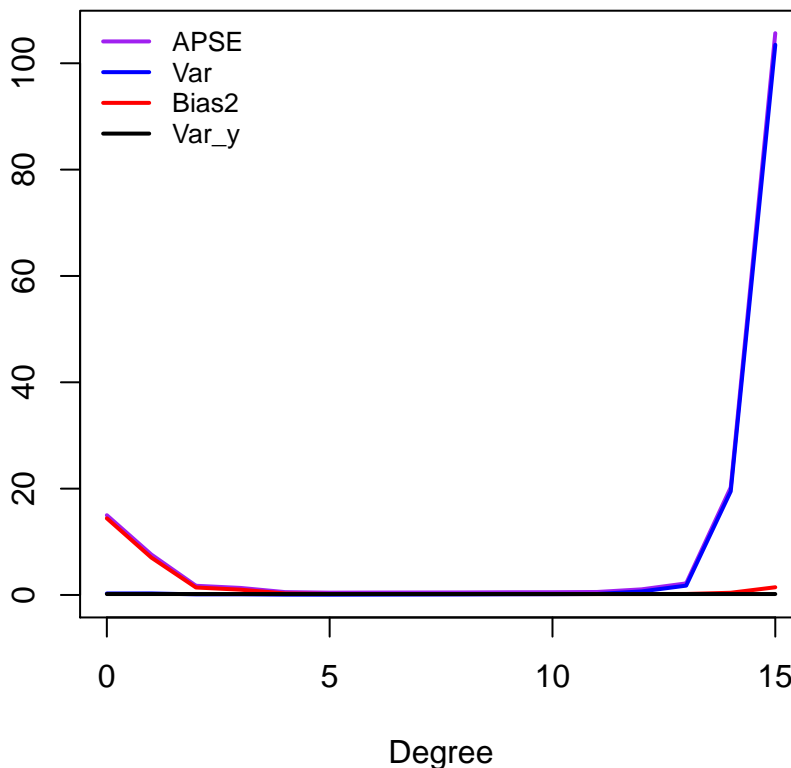
##      degrees      apse var_mutilde      bias2      var_y
## [1,]      0 15.00455 0.28249 14.42296 0.17746
## [2,]      1  7.59809 0.29122  7.02166 0.17746
## [3,]      2  1.73165 0.11673  1.40530 0.17746
## [4,]      3  1.32417 0.12218  0.99462 0.17746
## [5,]      4  0.53903 0.06021  0.28535 0.17746
## [6,]      5  0.43334 0.04975  0.19005 0.17746
## [7,]      6  0.45271 0.06744  0.18825 0.17746
## [8,]      7  0.46081 0.08152  0.17853 0.17746
## [9,]      8  0.47552 0.09539  0.17685 0.17746

```

```
## [10,]      9  0.50650      0.12853  0.17220 0.17746
## [11,]     10  0.52031      0.13963  0.17131 0.17746
## [12,]     11  0.56527      0.18261  0.17050 0.17746
## [13,]     12  1.05900      0.67066  0.17209 0.17746
## [14,]     13  2.17739      1.75961  0.19093 0.17746
## [15,]     14 20.23517     19.49681  0.41786 0.17746
## [16,]     15 105.66131    103.48264  1.44968 0.17746
```

*# Plot the results in a graph*


```
plot(degrees, apse_vals[1, ], xlab = "Degree", ylab = "", type = "l", ylim = c(0,
  max(apse_vals)), col = "purple", lwd = 2)
lines(degrees, apse_vals[2, ], col = "blue", lwd = 2)
lines(degrees, apse_vals[3, ], col = "red", lwd = 2)
lines(degrees, apse_vals[4, ], col = "black", lwd = 2)
legend("topleft", legend = c("APSE", "Var", "Bias2", "Var_y"), col = c("purple",
  "blue", "red", "black"), lwd = 2, bty = "n", cex = 0.8)
```



*# Make another 'zoomed-in' graph:*

```
plot(degrees, apse_vals[1, ], xlab = "Degree", ylab = "", type = "l", ylim = c(0,
  2), col = "purple", lwd = 2, main = "Zoomed In")
lines(degrees, apse_vals[2, ], col = "blue", lwd = 2)
lines(degrees, apse_vals[3, ], col = "red", lwd = 2)
lines(degrees, apse_vals[4, ], col = "black", lwd = 2)
legend("topright", legend = c("APSE", "Var", "Bias2", "Var_y"), col = c("purple",
  "blue", "red", "black"), lwd = 2, bty = "n", cex = 0.8)
```

## Zoomed In



(f) Plot the data again, and this time overlay the best polynomial as determined by your findings in (e).

```
# Plot the data
plot(fake, pch = 16, col = adjustcolor("darkgreen", 0.5))

# Get the best degree and fit a polynomial with this degree to all of
# the data.
best.deg <- degrees[which.min(apse_vals[1, ])]
print(best.deg)

## [1] 5

muhat.best <- getmuhat(fake, best.deg)

# Add the curve to the plot
curve(muhat.best, from = 0, to = 1, add = TRUE, col = "purple", lwd = 3)
```

